

STATISTICAL ANALYSIS

Cryptocurrency vs Equity Markets

MANAV AGARWAL

This notebook implements advanced statistical methods for comparing machine learning performance between cryptocurrency and equity markets, including deflated Sharpe ratios, probabilistic performance metrics, and robust hypothesis testing.

Data Flow:

- Input: `notebooks/01_comprehensive_results.pkl` from notebook 01
- Output: `notebooks/02_statistical_results.pkl` for downstream analysis

```
In [1]: # Import
import sys
import os
os.chdir('C:/Users/manav')
sys.path.append('src')

import pandas as pd
import numpy as np
from tqdm import tqdm # Progress bars
from joblib import Parallel, delayed # Parallel processing
import multiprocessing as mp
from functools import lru_cache
try:
    import cupy as cp
    HAS_GPU = True
    print("CuPy (GPU) available for acceleration")
except ImportError:
    HAS_GPU = False
    cp = None # Don't alias to np
    print("CuPy not available, using NumPy (CPU) fallback")

from datetime import datetime, timedelta
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from scipy.stats import jarque_bera, shapiro, anderson, kstest
from scipy.stats import skew, kurtosis, norm
from scipy.stats.mstats import normaltest
import warnings
warnings.filterwarnings('ignore')

from statsmodels.stats.diagnostic import acorr_ljungbox
from statsmodels.stats.stattools import jarque_bera as jb_test
from statsmodels.tsa.stattools import adfuller
from statsmodels.stats.diagnostic import het_arch
```

```

from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.ensemble import IsolationForest
import pickle
import json

plt.style.use('seaborn-v0_8-whitegrid')
sns.set_palette("husl")
plt.rcParams.update({
    'figure.dpi': 150,
    'figure.figsize': (14, 8),
    'font.size': 11,
    'axes.titleweight': 'bold',
    'axes.labelweight': 'bold',
    'axes.grid': True,
    'grid.alpha': 0.3,
    'lines.linewidth': 2
})

# Set number of parallel jobs
N_JOBS = mp.cpu_count() - 1 # Leave one CPU free
print(f"[OK] Libraries loaded successfully")
print(f"[INFO] Using {N_JOBS} parallel workers")

```

CuPy (GPU) available for acceleration
[OK] Libraries loaded successfully
[INFO] Using 31 parallel workers

1. Load Data from Previous Notebook

```

In [2]: # Load results from notebook 01
print("Loading data from notebook 01...")
try:
    with open('notebooks/01_comprehensive_results.pkl', 'rb') as f:
        nb01_results = pickle.load(f)
    print("Successfully loaded notebook 01 results")
    print(f"Keys available: {list(nb01_results.keys())}")
    all_data = nb01_results.get('all_data', {})
    crypto_symbols = nb01_results.get('crypto_symbols', [])
    equity_symbols = nb01_results.get('equity_symbols', [])
    metadata = nb01_results.get('metadata', {})
    engineered_data = {}
    model_results = {}
    config = metadata

    crypto_symbols = metadata.get('crypto_symbols', ['BTCUSD', 'ETHUSD', 'SOLUSD',
    equity_symbols = metadata.get('equity_symbols', ['SPY', 'QQQ', 'IWM'])
    print(f"\nCrypto symbols loaded: {crypto_symbols}")
    print(f"Equity symbols configured: {equity_symbols}")
    available_crypto = [s for s in crypto_symbols if s in all_data]
    available_equity = [s for s in equity_symbols if s in all_data]
    print(f"\nActually available:")
    print(f"  Crypto: {available_crypto} ({len(available_crypto)}/{len(crypto_symbols)})")
    print(f"  Equity: {available_equity} ({len(available_equity)}/{len(equity_symbols)})")

```

```
for symbol in available_crypto:
    if symbol in all_data:
        df = all_data[symbol]
        print(f"\n{symbol}:")
        print(f"  Records: {len(df)}")
        print(f"  Date range: {df.index.min()} to {df.index.max()}")
        print(f"  Columns: {list(df.columns)[:10]}") # Show first 10 columns

for symbol in available_equity:
    if symbol in all_data:
        df = all_data[symbol]
        print(f"\n{symbol}:")
        print(f"  Records: {len(df)}")
        print(f"  Date range: {df.index.min()} to {df.index.max()}")
        print(f"  Columns: {list(df.columns)[:10]}") # Show first 10 columns

except FileNotFoundError:
    print("Could not find Pickle file")
    raise
```

Loading data from notebook 01...

Successfully loaded notebook 01 results

Keys available: ['all_data', 'engineered_data', 'walk_forward_splits', 'model_results', 'statistical_results', 'config', 'timestamp']

Crypto symbols loaded: ['BTCUSD', 'ETHUSD', 'SOLUSD', 'XRPUSD', 'ADAUSD']

Equity symbols configured: ['SPY', 'QQQ', 'IWM']

Actually available:

Crypto: ['BTCUSD', 'ETHUSD', 'SOLUSD', 'XRPUSD', 'ADAUSD'] (5/5)

Equity: ['SPY', 'QQQ', 'IWM'] (3/3)

BTCUSD:

Records: 1359129

Date range: 2023-01-01 00:00:00 to 2025-08-01 23:59:00

Columns: ['ticker', 'volume', 'open', 'close', 'high', 'low', 'window_start', 'transactions']

ETHUSD:

Records: 1358218

Date range: 2023-01-01 00:00:00 to 2025-08-01 23:59:00

Columns: ['ticker', 'volume', 'open', 'close', 'high', 'low', 'window_start', 'transactions']

SOLUSD:

Records: 1354459

Date range: 2023-01-01 00:00:00 to 2025-08-01 23:59:00

Columns: ['ticker', 'volume', 'open', 'close', 'high', 'low', 'window_start', 'transactions']

XRPUSD:

Records: 1311371

Date range: 2023-01-01 00:00:00 to 2025-08-01 23:59:00

Columns: ['ticker', 'volume', 'open', 'close', 'high', 'low', 'window_start', 'transactions']

ADAUSD:

Records: 1347195

Date range: 2023-01-01 00:00:00 to 2025-08-01 23:59:00

Columns: ['ticker', 'volume', 'open', 'close', 'high', 'low', 'window_start', 'transactions']

SPY:

Records: 77639

Date range: 2023-01-03 09:00:00 to 2024-01-13 00:59:00

Columns: ['ticker', 'volume', 'open', 'close', 'high', 'low', 'window_start', 'transactions']

QQQ:

Records: 78493

Date range: 2023-01-03 09:00:00 to 2024-01-13 00:59:00

Columns: ['ticker', 'volume', 'open', 'close', 'high', 'low', 'window_start', 'transactions']

IWM:

Records: 66382

Date range: 2023-01-03 09:00:00 to 2024-01-13 00:59:00
Columns: ['ticker', 'volume', 'open', 'close', 'high', 'low', 'window_start', 'transactions']

2. Data Quality Assessment and Preparation

```
In [3]: # Prepare data for analysis
analysis_data = {}
# Process available symbols
all_symbols = available_crypto + available_equity

print("Preparing data for statistical analysis...\n")

for symbol in all_symbols:
    if symbol in all_data:
        df = all_data[symbol].copy()

        # Calculate returns if not present
        if 'returns' not in df.columns:
            df['returns'] = df['close'].pct_change()

        # Calculate log returns
        if 'log_returns' not in df.columns:
            df['log_returns'] = np.log(df['close'] / df['close'].shift(1))

        df = df.dropna()

        analysis_data[symbol] = df

    # Basic statistics
    print(f"{symbol}:")
    print(f"  Shape: {df.shape}")
    print(f"  Missing values: {df.isnull().sum().sum()}")
    print(f"  Return mean: {df['returns'].mean():.6f}")
    print(f"  Return std: {df['returns'].std():.6f}")
    print(f"  Skewness: {skew(df['returns']):.3f}")
    print(f"  Kurtosis: {kurtosis(df['returns']):.3f}")
    print()

print(f"\n[OK] Prepared {len(analysis_data)} datasets for analysis")
```

Preparing data for statistical analysis...

BTCUSD:

Shape: (1359128, 10)
Missing values: 0
Return mean: 0.000002
Return std: 0.000810
Skewness: -0.291
Kurtosis: 44.698

ETHUSD:

Shape: (1358217, 10)
Missing values: 0
Return mean: 0.000001
Return std: 0.000925
Skewness: -2.114
Kurtosis: 293.596

SOLUSD:

Shape: (1354458, 10)
Missing values: 0
Return mean: 0.000003
Return std: 0.001401
Skewness: 1.755
Kurtosis: 301.689

XRPUSD:

Shape: (1311370, 10)
Missing values: 0
Return mean: 0.000003
Return std: 0.001355
Skewness: -0.663
Kurtosis: 270.586

ADAUSD:

Shape: (1347194, 10)
Missing values: 0
Return mean: 0.000002
Return std: 0.001416
Skewness: 0.806
Kurtosis: 542.085

SPY:

Shape: (77638, 10)
Missing values: 0
Return mean: 0.000003
Return std: 0.000608
Skewness: 135.796
Kurtosis: 29773.755

QQQ:

Shape: (78492, 10)
Missing values: 0
Return mean: 0.000006
Return std: 0.000944
Skewness: 171.214

Kurtosis: 40704.727

IWM:

Shape: (66381, 10)
Missing values: 0
Return mean: 0.000002
Return std: 0.000713
Skewness: 54.524
Kurtosis: 8408.172

[OK] Prepared 8 datasets for analysis

3. Distribution Analysis and Normality Testing

Comprehensive distribution analysis for returns

Formulas:

- Skewness: $E[(X - \mu)^3] / \sigma^3$
- Kurtosis: $E[(X - \mu)^4] / \sigma^4$
- Jarque-Bera: $n/6 * (S^2 + (K-3)^2/4)$ where S =skewness, K =kurtosis
- Sharpe Ratio: $(E[R] - R_f) / \sigma$, we assume $R_f = 0$

```
In [4]: def analyze_distribution(data, symbol, returns_col='returns'):
    try:
        returns = data[returns_col].dropna()
        returns_np = returns.values if hasattr(returns, 'values') else np.array(returns)

        # Calculate moments
        mean_return = float(np.mean(returns_np))
        std_return = float(np.std(returns_np, ddof=1))
        moments = {
            'mean': mean_return,
            'std': std_return,
            'skewness': float(skew(returns_np, bias=False)), # Unbiased estimator
            'kurtosis': float(kurtosis(returns_np, bias=False)), # Unbiased estimator
            'excess_kurtosis': float(kurtosis(returns_np, fisher=True, bias=False))
        }

        # Normality tests
        jb_stat, jb_pval = jarque_bera(returns_np)
        shapiro_stat, shapiro_pval = shapiro(returns_np) if len(returns_np) < 5000 else None
        ks_stat, ks_pval = kstest(returns_np, 'norm', args=(mean_return, std_return))

        # ARCH effect test
        try:
            arch_test = het_arch(returns_np, nlags=10)
            has_arch = arch_test[1] < 0.05
        except:
            has_arch = None

        # Dickey-Fuller test for stationarity
        try:
            adf_result = adfuller(returns_np, autolag='AIC')
```

```

        adf_stat = adf_result[0]
        adf_pval = adf_result[1]
        is_stationary = adf_pval < 0.05
    except:
        adf_stat, adf_pval, is_stationary = None, None, None

    results = {
        'symbol': symbol,
        'moments': moments,
        'normality': {
            'jarque_bera': {'statistic': jb_stat, 'p_value': jb_pval},
            'shapiro': {'statistic': shapiro_stat, 'p_value': shapiro_pval},
            'ks_test': {'statistic': ks_stat, 'p_value': ks_pval}
        },
        'arch_effects': has_arch,
        'stationarity': {
            'adf_statistic': adf_stat,
            'p_value': adf_pval,
            'is_stationary': is_stationary
        }
    }

    return results
except Exception as e:
    print(f"Error analyzing {symbol}: {str(e)}")
    return None

print("DISTRIBUTION ANALYSIS RESULTS")
print("="*80)
print(f"Processing {len(all_symbols)} symbols\n")

distribution_results = {}

# Process each symbol individually with progress tracking
for i, symbol in enumerate(all_symbols, 1):
    if symbol in analysis_data:
        print(f"[{i}/{len(all_symbols)}] Processing {symbol}...", end=' ')
        result = analyze_distribution(analysis_data[symbol], symbol)
        if result:
            distribution_results[symbol] = result
            print("OK")
        else:
            print("FAIL")
    else:
        print(f"[{i}/{len(all_symbols)}] Skipping {symbol} (no data)")

# Display results
print("RESULTS SUMMARY")
print("="*80)

for symbol in all_symbols:
    if symbol in distribution_results:
        results = distribution_results[symbol]
        print(f"\n{symbol}:")
        print(f"  Moments:")
        for k, v in results['moments'].items():

```



```

        print(f"    {k}: {v:.6f}")
    print(f" Normality (p-values):")
    for test, vals in results['normality'].items():
        if vals['p_value'] is not None and not np.isnan(vals['p_value']):
            print(f"    {test}: {vals['p_value']:.6f} {'[NORMAL]' if vals['p_value'] < 0.05 else '[NOT NORMAL]'}")
    print(f" ARCH effects: {'Yes' if results['arch_effects'] else 'No' if results['arch_effects'] is not None else 'None'}")
    if results['stationarity']['is_stationary'] is not None:
        print(f" Stationary: {'Yes' if results['stationarity']['is_stationary'] else 'No' if results['stationarity']['is_stationary'] is not None else 'None'}")

print(f"\n[OK] Distribution analysis complete for {len(distribution_results)} symbols")

```

DISTRIBUTION ANALYSIS RESULTS

Processing 8 symbols sequentially

[1/8] Processing BTCUSD... ✓
[2/8] Processing ETHUSD... ✓
[3/8] Processing SOLUSD... ✓
[4/8] Processing XRPUSD... ✓
[5/8] Processing ADAUSD... ✓
[6/8] Processing SPY... ✓
[7/8] Processing QQQ... ✓
[8/8] Processing IWM... ✓

RESULTS SUMMARY

BTCUSD:

Moments:

mean: 0.000002
std: 0.000810
skewness: -0.291128
kurtosis: 44.698513
excess_kurtosis: 44.698513

Normality (p-values):

jarque_bera: 0.000000 [NON-NORMAL]
ks_test: 0.000000 [NON-NORMAL]

ARCH effects: Yes

ETHUSD:

Moments:

mean: 0.000001
std: 0.000925
skewness: -2.114286
kurtosis: 293.597271
excess_kurtosis: 293.597271

Normality (p-values):

jarque_bera: 0.000000 [NON-NORMAL]
ks_test: 0.000000 [NON-NORMAL]

ARCH effects: Yes

SOLUSD:

Moments:

mean: 0.000003
std: 0.001401
skewness: 1.754651
kurtosis: 301.690421
excess_kurtosis: 301.690421

Normality (p-values):

jarque_bera: 0.000000 [NON-NORMAL]
ks_test: 0.000000 [NON-NORMAL]

ARCH effects: Yes

XRPUSD:

Moments:

mean: 0.000003

std: 0.001355
skewness: -0.662806
kurtosis: 270.587504
excess_kurtosis: 270.587504
Normality (p-values):
jarque_bera: 0.000000 [NON-NORMAL]
ks_test: 0.000000 [NON-NORMAL]
ARCH effects: Yes

ADAUSD:

Moments:
mean: 0.000002
std: 0.001416
skewness: 0.805913
kurtosis: 542.086541
excess_kurtosis: 542.086541
Normality (p-values):
jarque_bera: 0.000000 [NON-NORMAL]
ks_test: 0.000000 [NON-NORMAL]
ARCH effects: Yes

SPY:

Moments:
mean: 0.000003
std: 0.000608
skewness: 135.798398
kurtosis: 29775.672142
excess_kurtosis: 29775.672142
Normality (p-values):
jarque_bera: 0.000000 [NON-NORMAL]
ks_test: 0.000000 [NON-NORMAL]
ARCH effects: No
Stationary: Yes

QQQ:

Moments:
mean: 0.000006
std: 0.000944
skewness: 171.216973
kurtosis: 40707.320176
excess_kurtosis: 40707.320176
Normality (p-values):
jarque_bera: 0.000000 [NON-NORMAL]
ks_test: 0.000000 [NON-NORMAL]
ARCH effects: No
Stationary: Yes

IWM:

Moments:
mean: 0.000002
std: 0.000713
skewness: 54.524786
kurtosis: 8408.805944
excess_kurtosis: 8408.805944
Normality (p-values):
jarque_bera: 0.000000 [NON-NORMAL]

ks_test: 0.000000 [NON-NORMAL]

ARCH effects: No

Stationary: Yes

[OK] Distribution analysis complete for 8 symbols

4. Comprehensive Return Distribution Visualization

```
In [5]: # Visualize distributions for all available symbols
print("DISTRIBUTION VISUALIZATION")
print("="*80)

if all_symbols:
    n_symbols = len(all_symbols)
    n_cols = min(3, n_symbols)
    n_rows = (n_symbols + n_cols - 1) // n_cols

    # Print distribution statistics first
    print("\nDistribution Statistics Summary:")
    for symbol in all_symbols:
        if symbol in analysis_data:
            returns = analysis_data[symbol]['returns'].dropna()

            print(f"\n{symbol}:")
            print(f"  Count: {len(returns)}")
            print(f"  Mean: {returns.mean():.6f}")
            print(f"  Std: {returns.std():.6f}")
            print(f"  Skewness: {skew(returns):.3f}")
            print(f"  Kurtosis: {kurtosis(returns):.3f}")
            print(f"  Min: {returns.min():.6f}")
            print(f"  25%: {returns.quantile(0.25):.6f}")
            print(f"  50% (median): {returns.quantile(0.50):.6f}")
            print(f"  75%: {returns.quantile(0.75):.6f}")
            print(f"  Max: {returns.max():.6f}")

    # Create visualizations only if we have symbols
    if n_symbols > 0:
        try:
            fig, axes = plt.subplots(n_rows, n_cols, figsize=(15, 4*n_rows))
            if n_symbols == 1:
                axes = [axes]
            else:
                axes = axes.flatten() if n_rows > 1 else axes

            for idx, symbol in enumerate(all_symbols):
                if symbol in analysis_data:
                    ax = axes[idx]
                    returns = analysis_data[symbol]['returns'].dropna()

                    # Plot histogram
                    n, bins, patches = ax.hist(returns, bins=50, density=True, alpha=0.7,
                                                color='blue', edgecolor='black')

                    # Fit normal distribution
                    mu, sigma = returns.mean(), returns.std()
```

```

x = np.linspace(returns.min(), returns.max(), 100)
ax.plot(x, norm.pdf(x, mu, sigma), 'r-', linewidth=2, label='No

# Add statistics
stats_text = f'Mean: {mu:.4f}\nStd: {sigma:.4f}\nSkew: {skew(re
ax.text(0.02, 0.98, stats_text, transform=ax.transAxes,
        fontsize=9, verticalalignment='top',
        bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5

ax.set_title(f'{symbol} Return Distribution', fontweight='bold'
ax.set_xlabel('Returns')
ax.set_ylabel('Density')
ax.legend()
ax.grid(True, alpha=0.3)

# Remove empty subplots
for idx in range(n_symbols, len(axes)):
    fig.delaxes(axes[idx])

plt.tight_layout()
plt.suptitle('Return Distributions Across Assets', fontsize=16, fontwei
plt.show()
except Exception as e:
    print(f"Could not create distribution plots: {e}")
else:
    print("No symbols available for distribution visualization")

```

DISTRIBUTION VISUALIZATION

=====

Distribution Statistics Summary:

BTCUSD:

Count: 1359128
Mean: 0.000002
Std: 0.000810
Skewness: -0.291
Kurtosis: 44.698
Min: -0.036938
25%: -0.000286
50% (median): 0.000000
75%: 0.000290
Max: 0.036828

ETHUSD:

Count: 1358217
Mean: 0.000001
Std: 0.000925
Skewness: -2.114
Kurtosis: 293.596
Min: -0.102181
25%: -0.000351
50% (median): 0.000000
75%: 0.000353
Max: 0.054582

SOLUSD:

Count: 1354458
Mean: 0.000003
Std: 0.001401
Skewness: 1.755
Kurtosis: 301.689
Min: -0.059722
25%: -0.000602
50% (median): 0.000000
75%: 0.000601
Max: 0.190239

XRPUSD:

Count: 1311370
Mean: 0.000003
Std: 0.001355
Skewness: -0.663
Kurtosis: 270.586
Min: -0.092770
25%: -0.000485
50% (median): 0.000000
75%: 0.000490
Max: 0.113176

ADAUSD:

Count: 1347194
Mean: 0.000002

Std: 0.001416
Skewness: 0.806
Kurtosis: 542.085
Min: -0.138424
25%: -0.000560
50% (median): 0.000000
75%: 0.000562
Max: 0.169852

SPY:

Count: 77638
Mean: 0.000003
Std: 0.000608
Skewness: 135.796
Kurtosis: 29773.755
Min: -0.012598
25%: -0.000145
50% (median): 0.000000
75%: 0.000147
Max: 0.133223

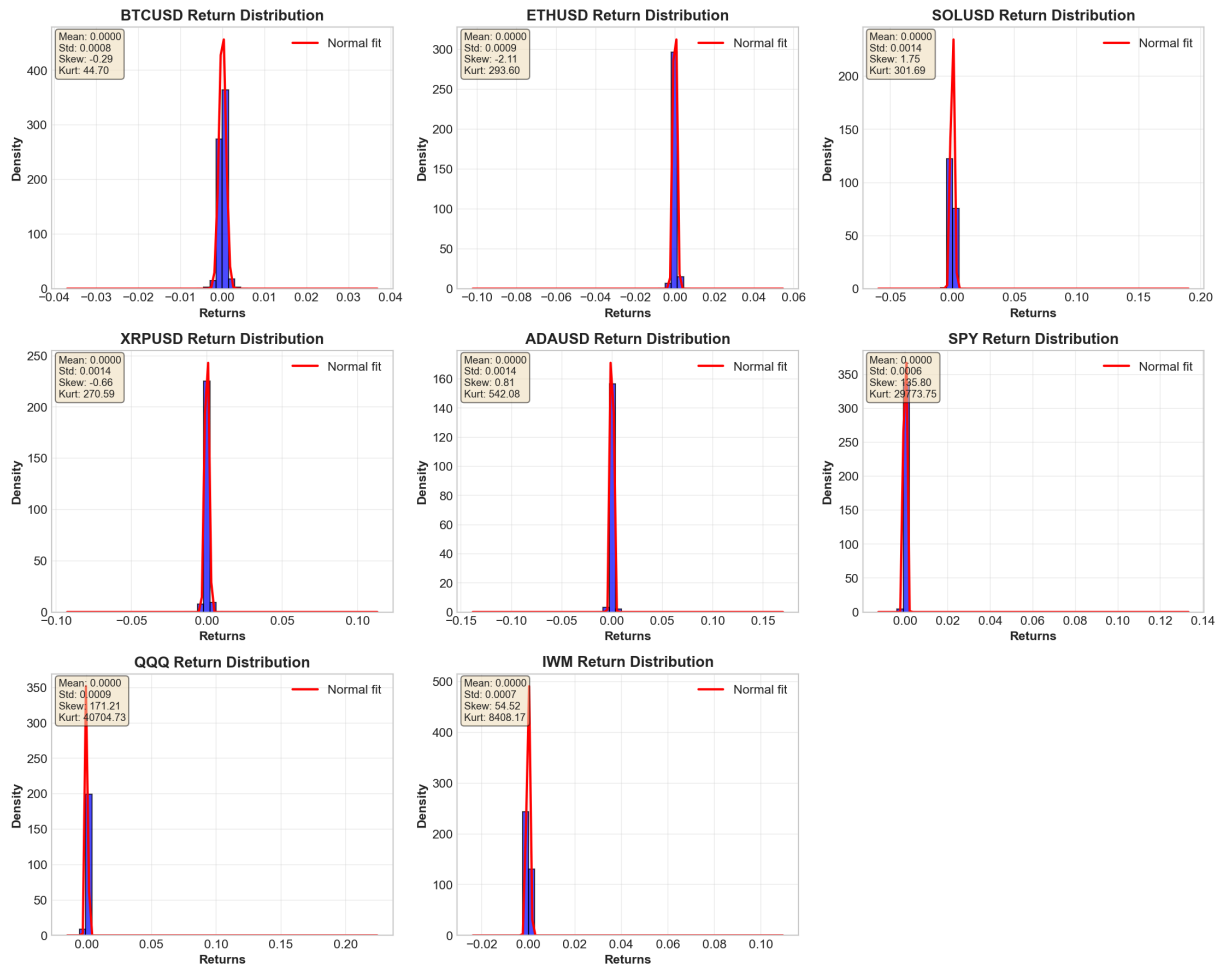
QQQ:

Count: 78492
Mean: 0.000006
Std: 0.000944
Skewness: 171.214
Kurtosis: 40704.727
Min: -0.015085
25%: -0.000186
50% (median): 0.000000
75%: 0.000190
Max: 0.224497

IWM:

Count: 66381
Mean: 0.000002
Std: 0.000713
Skewness: 54.524
Kurtosis: 8408.172
Min: -0.023872
25%: -0.000233
50% (median): 0.000000
75%: 0.000252
Max: 0.109438

Return Distributions Across Assets



[OK] Distribution visualizations complete

5. Correlation Analysis

```
In [6]: # Optimized Correlation Analysis with vectorization
print("CORRELATION ANALYSIS")
print("="*80)

print("Building returns matrix...")

# Get all returns data at once
returns_dict = {}
for symbol in all_symbols:
    if symbol in analysis_data:
        returns_dict[symbol] = analysis_data[symbol]['returns']
        print(f" Added {symbol}: {len(analysis_data[symbol]['returns'])} data points")

if returns_dict:
    returns_matrix = pd.DataFrame(returns_dict)

    # Align all series to common dates (removes NaN)
    initial_shape = returns_matrix.shape
    returns_matrix = returns_matrix.dropna()
    final_shape = returns_matrix.shape
```



```

print(f"\nReturns matrix shape: {initial_shape} -> {final_shape} after removing
if len(returns_matrix) > 0:
    print(f>Date range: {returns_matrix.index.min()} to {returns_matrix.index.m
    print(f"Common data points: {len(returns_matrix)}")
    print(f"Symbols in matrix: {list(returns_matrix.columns)}")
else:
    print("WARNING: No common data points found across symbols!")
    returns_matrix = pd.DataFrame()
else:
    returns_matrix = pd.DataFrame()
    print("No returns data available")

if len(returns_matrix.columns) > 1 and len(returns_matrix) > 0:
    print("\nCalculating correlation matrix...")
    correlation_matrix = returns_matrix.corr(method='pearson') # Pearson correlati

    spearman_corr = returns_matrix.corr(method='spearman')

    print("\nPearson Correlation Matrix:")
    print(correlation_matrix.round(3))

    print("\nSpearman Correlation Matrix:")
    print(spearman_corr.round(3))

    print("\nKey Correlations (Pearson):")
    for i, sym1 in enumerate(correlation_matrix.columns):
        for j, sym2 in enumerate(correlation_matrix.columns):
            if i < j: # Only print upper triangle
                corr_val = correlation_matrix.loc[sym1, sym2]
                print(f" {sym1} vs {sym2}: {corr_val:.3f}")

    try:
        fig, axes = plt.subplots(1, 2, figsize=(16, 6))

        # Pearson correlation
        sns.heatmap(correlation_matrix, annot=True, fmt='.2f', cmap='coolwarm', cen
                    vmin=-1, vmax=1, square=True, linewidths=1,
                    cbar_kws={"shrink": 0.8}, ax=axes[0])
        axes[0].set_title('Pearson Correlation Matrix', fontsize=14, fontweight='bo

        # Spearman correlation
        sns.heatmap(spearman_corr, annot=True, fmt='.2f', cmap='coolwarm', center=0
                    vmin=-1, vmax=1, square=True, linewidths=1,
                    cbar_kws={"shrink": 0.8}, ax=axes[1])
        axes[1].set_title('Spearman Rank Correlation Matrix', fontsize=14, fontweig

        plt.tight_layout()
        plt.show()
    except Exception as e:
        print(f"Could not create correlation heatmap: {e}")
else:
    correlation_matrix = pd.DataFrame()
    print("\nInsufficient data")

if len(returns_matrix.columns) >= 2 and len(returns_matrix) > 0:
    symbol1, symbol2 = returns_matrix.columns[0], returns_matrix.columns[1]

```

```

window = min(30 * 24, len(returns_matrix) // 2) # 30 days or half the data

print(f"\n" + "="*80)
print(f"ROLLING CORRELATION ANALYSIS")
print(f"Symbols: {symbol1} vs {symbol2}")
print(f"Window: {window} hours (~{window/24:.1f} days)")
print("="*80)

try:
    subset_data = returns_matrix[[symbol1, symbol2]]
    rolling_corr_series = subset_data[symbol1].rolling(
        window=window, min_periods=max(10, window//4)
    ).corr(subset_data[symbol2])

    rolling_corr_series = rolling_corr_series.dropna()

    if len(rolling_corr_series) > 0:
        # Calculate correlation stability metrics
        corr_mean = rolling_corr_series.mean()
        corr_std = rolling_corr_series.std()
        corr_min = rolling_corr_series.min()
        corr_max = rolling_corr_series.max()
        corr_median = rolling_corr_series.median()

        print(f"\nRolling Correlation Statistics:")
        print(f"  Data points: {len(rolling_corr_series)}")
        print(f"  Mean: {corr_mean:.4f}")
        print(f"  Median: {corr_median:.4f}")
        print(f"  Std: {corr_std:.4f}")
        print(f"  Min: {corr_min:.4f}")
        print(f"  Max: {corr_max:.4f}")
        print(f"  Range: {corr_max - corr_min:.4f}")

        print(f"\nSample correlation values (first 10):")
        for idx, (date, val) in enumerate(rolling_corr_series.head(10).items()):
            print(f"  {date}: {val:.4f}")

        print(f"\nSample correlation values (last 10):")
        for idx, (date, val) in enumerate(rolling_corr_series.tail(10).items()):
            print(f"  {date}: {val:.4f}")

        # Correlation stability analysis
        print(f"\nCorrelation Stability Analysis:")
        high_corr_periods = (rolling_corr_series > 0.7).sum()
        low_corr_periods = (rolling_corr_series < 0.3).sum()
        negative_corr_periods = (rolling_corr_series < 0).sum()

        print(f"Periods with correlation > 0.7: {high_corr_periods} ({high_corr_p")
        print(f"Periods with correlation < 0.3: {low_corr_periods} ({low_corr_p")
        print(f"Periods with negative correlation: {negative_corr_periods} ({ne

try:
    plt.figure(figsize=(14, 6))
    plt.plot(rolling_corr_series.index, rolling_corr_series.values,
             linewidth=1, alpha=0.8, label='Rolling Correlation')
    plt.axhline(y=corr_mean, color='r', linestyle='--',

```

```

        label=f'Mean: {corr_mean:.3f}', linewidth=2)
plt.axhline(y=corr_median, color='g', linestyle='--',
            label=f'Median: {corr_median:.3f}', linewidth=1, alpha=0.2)
plt.fill_between(rolling_corr_series.index,
                 corr_mean - corr_std,
                 corr_mean + corr_std,
                 alpha=0.2, color='gray', label=f'±1 Std: {corr_std:.3f}')

plt.title(f'{window}-Hour Rolling Correlation: {symbol1} vs {symbol2}')
plt.xlabel('Date')
plt.ylabel('Correlation')
plt.ylim(-1, 1)
plt.legend(loc='best')
plt.grid(True, alpha=0.3)

# Add text box with statistics
textstr = f'Points: {len(rolling_corr_series)}\nMin: {corr_min:.3f}\nMax: {corr_max:.3f}'
plt.text(0.02, 0.95, textstr, transform=plt.gca().transAxes,
        fontsize=10, verticalalignment='top',
        bbox=dict(boxstyle='round', facecolor='wheat', alpha=0.5))

plt.tight_layout()
plt.show()
except Exception as e:
    print(f"Could not create rolling correlation plot: {e}")

else:
    print("No valid rolling correlation data points")

except Exception as e:
    print(f"Error calculating rolling correlation: {str(e)}")
import traceback
traceback.print_exc()

```

CORRELATION ANALYSIS

=====

Building returns matrix...

Added BTCUSD: 1359128 data points
Added ETHUSD: 1358217 data points
Added SOLUSD: 1354458 data points
Added XRPUSD: 1311370 data points
Added ADAUSD: 1347194 data points
Added SPY: 77638 data points
Added QQQ: 78492 data points
Added IWM: 66381 data points

Returns matrix shape: (1359359, 8) -> (54954, 8) after removing NaN

Date range: 2023-01-03 09:01:00 to 2024-01-13 00:59:00

Common data points: 54954

Symbols in matrix: ['BTCUSD', 'ETHUSD', 'SOLUSD', 'XRPUSD', 'ADAUSD', 'SPY', 'QQQ', 'IWM']

Calculating correlation matrix...

Pearson Correlation Matrix:

	BTCUSD	ETHUSD	SOLUSD	XRPUSD	ADAUSD	SPY	QQQ	IWM
BTCUSD	1.000	0.749	0.642	0.472	0.606	0.197	0.163	0.220
ETHUSD	0.749	1.000	0.671	0.506	0.641	0.209	0.174	0.229
SOLUSD	0.642	0.671	1.000	0.480	0.623	0.194	0.162	0.213
XRPUSD	0.472	0.506	0.480	1.000	0.489	0.128	0.106	0.146
ADAUSD	0.606	0.641	0.623	0.489	1.000	0.174	0.147	0.192
SPY	0.197	0.209	0.194	0.128	0.174	1.000	0.944	0.835
QQQ	0.163	0.174	0.162	0.106	0.147	0.944	1.000	0.789
IWM	0.220	0.229	0.213	0.146	0.192	0.835	0.789	1.000



Spearman Correlation Matrix:

	BTCUSD	ETHUSD	SOLUSD	XRPUSD	ADAUSD	SPY	QQQ	IWM
BTCUSD	1.000	0.757	0.641	0.430	0.589	0.377	0.379	0.323
ETHUSD	0.757	1.000	0.661	0.455	0.620	0.384	0.385	0.325
SOLUSD	0.641	0.661	1.000	0.433	0.587	0.329	0.330	0.284
XRPUSD	0.430	0.455	0.433	1.000	0.431	0.214	0.215	0.189
ADAUSD	0.589	0.620	0.587	0.431	1.000	0.292	0.297	0.252
SPY	0.377	0.384	0.329	0.214	0.292	1.000	0.864	0.739
QQQ	0.379	0.385	0.330	0.215	0.297	0.864	1.000	0.659
IWM	0.323	0.325	0.284	0.189	0.252	0.739	0.659	1.000

Key Correlations (Pearson):

BTCUSD vs ETHUSD: 0.749
BTCUSD vs SOLUSD: 0.642
BTCUSD vs XRPUSD: 0.472
BTCUSD vs ADAUSD: 0.606
BTCUSD vs SPY: 0.197
BTCUSD vs QQQ: 0.163
BTCUSD vs IWM: 0.220
ETHUSD vs SOLUSD: 0.671
ETHUSD vs XRPUSD: 0.506
ETHUSD vs ADAUSD: 0.641
ETHUSD vs SPY: 0.209
ETHUSD vs QQQ: 0.174
ETHUSD vs IWM: 0.229

SOLUSD vs XRPUSD: 0.480
SOLUSD vs ADAUSD: 0.623
SOLUSD vs SPY: 0.194
SOLUSD vs QQQ: 0.162
SOLUSD vs IWM: 0.213
XRPUSD vs ADAUSD: 0.489
XRPUSD vs SPY: 0.128
XRPUSD vs QQQ: 0.106
XRPUSD vs IWM: 0.146
ADAUSD vs SPY: 0.174
ADAUSD vs QQQ: 0.147
ADAUSD vs IWM: 0.192
SPY vs QQQ: 0.944
SPY vs IWM: 0.835
QQQ vs IWM: 0.789



```
=====
ROLLING CORRELATION ANALYSIS
Symbols: BTCUSD vs ETHUSD
Window: 720 hours (~30.0 days)
=====
```

Rolling Correlation Statistics:

```
Data points: 54775
Mean: 0.7754
Median: 0.8169
Std: 0.1268
Min: 0.1095
Max: 0.9240
Range: 0.8146
```

Sample correlation values (first 10):


```
2023-01-03 14:25:00: 0.1095
2023-01-03 14:27:00: 0.1095
2023-01-03 14:28:00: 0.1102
2023-01-03 14:30:00: 0.1165
2023-01-03 14:31:00: 0.1168
2023-01-03 14:32:00: 0.1254
2023-01-03 14:33:00: 0.1289
2023-01-03 14:34:00: 0.1356
2023-01-03 14:35:00: 0.1340
2023-01-03 14:36:00: 0.1469
```

Sample correlation values (last 10):

```
2024-01-12 22:30:00: 0.7130
2024-01-12 22:33:00: 0.7136
2024-01-12 22:34:00: 0.7143
2024-01-12 23:07:00: 0.7142
2024-01-12 23:08:00: 0.7154
2024-01-12 23:12:00: 0.7154
2024-01-13 00:56:00: 0.7153
2024-01-13 00:57:00: 0.7153
2024-01-13 00:58:00: 0.7154
2024-01-13 00:59:00: 0.7150
```

Correlation Stability Analysis:

```
Periods with correlation > 0.7: 45170 (82.5%)
Periods with correlation < 0.3: 689 (1.3%)
Periods with negative correlation: 0 (0.0%)
```



[OK] Correlation analysis complete

6. Volatility Analysis and GARCH Effects

```
In [19]: print("VOLATILITY ANALYSIS")
print("="*80)

def calculate_volatility_metrics(symbol, df):
    try:
        returns = df['returns'].dropna()

        # Basic volatility metrics
        daily_vol = returns.std()
        annualization_factor = np.sqrt(365 * 24) # For hourly data

        metrics = {
            'daily_vol': daily_vol,
            'annualized_vol': daily_vol * annualization_factor,
            'vol_of_vol': returns.rolling(window=24).std().std(),
            'max_drawdown': (df['close'] / df['close'].cummax() - 1).min(),
            'realized_vol_30d': returns.rolling(window=24*30).std().mean() if len(r
            'realized_vol_7d': returns.rolling(window=24*7).std().mean() if len(ret
            'sharpe_ratio': (returns.mean() / daily_vol) * annualization_factor if
        }

        # GARCH test for volatility
        try:
            arch_test = het_arch(returns.values, nlags=10)
            metrics['arch_lm_stat'] = arch_test[0]
            metrics['arch_lm_pvalue'] = arch_test[1]
            metrics['has_volatility_clustering'] = arch_test[1] < 0.05
        except:
            metrics['arch_lm_stat'] = None
            metrics['arch_lm_pvalue'] = None
            metrics['has_volatility_clustering'] = None

        return symbol, metrics
    except Exception as e:
        print(f"Error calculating volatility for {symbol}: {str(e)}")
        return symbol, {}
```

```

print(f"Calculating volatility metrics for {len(all_symbols)} symbols...\n")

volatility_results = {}

for i, symbol in enumerate(all_symbols, 1):
    if symbol in analysis_data:
        print(f"[{i}/{len(all_symbols)}] Processing {symbol}...", end=' ')
        sym, metrics = calculate_volatility_metrics(symbol, analysis_data[symbol])
        if metrics:
            volatility_results[sym] = metrics
            print("OK")
        else:
            print("FAIL")

if volatility_results:
    vol_df = pd.DataFrame(volatility_results).T

    print("\nVolatility Metrics Summary:")
    print("="*80)
    print(vol_df.round(4))

    # Statistical comparison between asset classes
    if len(available_crypto) > 0 and len(available_equity) > 0:
        crypto_vols = [volatility_results[s]['annualized_vol'] for s in available_crypto]
        equity_vols = [volatility_results[s]['annualized_vol'] for s in available_equity]

        if crypto_vols and equity_vols:
            crypto_vol = np.mean(crypto_vols)
            equity_vol = np.mean(equity_vols)

            print(f"\nVolatility Comparison:")
            print(f"  Crypto avg annualized vol: {crypto_vol:.2%}")
            print(f"  Equity avg annualized vol: {equity_vol:.2%}")
            print(f"  Ratio (Crypto/Equity): {crypto_vol/equity_vol:.2f}x")

    if len(volatility_results) > 1:
        fig, axes = plt.subplots(1, 3, figsize=(18, 6))

        # 1. Annualized volatility comparison
        ax = axes[0]
        symbols = list(volatility_results.keys())
        vols = [volatility_results[s].get('annualized_vol', 0) for s in symbols]
        colors = ['blue' if s in available_crypto else 'green' for s in symbols]
        bars = ax.bar(range(len(symbols)), vols, color=colors, alpha=0.7)
        ax.set_xticks(range(len(symbols)))
        ax.set_xticklabels(symbols, rotation=45)
        ax.set_title('Annualized Volatility Comparison', fontweight='bold')
        ax.set_ylabel('Volatility')
        ax.grid(True, alpha=0.3)

        # 2. Sharpe ratio comparison
        ax = axes[1]
        sharpes = [volatility_results[s].get('sharpe_ratio', 0) for s in symbols]

```



```

bars = ax.bar(range(len(symbols)), sharpes, color=colors, alpha=0.7)
ax.set_xticks(range(len(symbols)))
ax.set_xticklabels(symbols, rotation=45)
ax.set_title('Sharpe Ratio Comparison', fontweight='bold')
ax.set_ylabel('Sharpe Ratio')
ax.axhline(y=0, color='r', linestyle='--', alpha=0.5)
ax.grid(True, alpha=0.3)

# 3. Max drawdown
ax = axes[2]
drawdowns = [abs(volatility_results[s].get('max_drawdown', 0)) * 100 for s in symbols]
bars = ax.bar(range(len(symbols)), drawdowns, color=colors, alpha=0.7)
ax.set_xticks(range(len(symbols)))
ax.set_xticklabels(symbols, rotation=45)
ax.set_title('Maximum Drawdown', fontweight='bold')
ax.set_ylabel('Drawdown (%)')
ax.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

```

VOLATILITY ANALYSIS

Calculating volatility metrics for 8 symbols...

[1/8] Processing BTCUSD... OK
[2/8] Processing ETHUSD... OK
[3/8] Processing SOLUSD... OK
[4/8] Processing XRPUSD... OK
[5/8] Processing ADAUSD... OK
[6/8] Processing SPY... OK
[7/8] Processing QQQ... OK
[8/8] Processing IWM... OK

Volatility Metrics Summary:


	daily_vol	annualized_vol	vol_of_vol	max_drawdown	realized_vol_30d \
BTCUSD	0.00081	0.075781	0.000483	-0.33124	0.000717
ETHUSD	0.000925	0.086596	0.000562	-0.661692	0.000814
SOLUSD	0.001401	0.131101	0.000767	-0.675678	0.001265
XRPUSD	0.001355	0.126818	0.00089	-0.587128	0.001128
ADAUSD	0.001416	0.132569	0.00087	-0.656486	0.001213
SPY	0.000608	0.056869	0.000522	-0.089485	0.000397
QQQ	0.000944	0.088379	0.000851	-0.090068	0.000546
IWM	0.000713	0.066738	0.000523	-0.158912	0.000581

	realized_vol_7d	sharpe_ratio	arch_lm_stat	arch_lm_pvalue \
BTCUSD	0.000691	0.201567	157132.30428	0.0
ETHUSD	0.000781	0.12309	160925.405202	0.0
SOLUSD	0.001225	0.203151	15657.134009	0.0
XRPUSD	0.001082	0.177495	191599.083528	0.0
ADAUSD	0.001173	0.118587	280458.703361	0.0
SPY	0.000356	0.442192	0.000636	1.0
QQQ	0.000478	0.569295	0.000268	1.0
IWM	0.000538	0.20791	0.43298	0.999997

	has_volatility_clustering
BTCUSD	True
ETHUSD	True
SOLUSD	True
XRPUSD	True
ADAUSD	True
SPY	False
QQQ	False
IWM	False

Volatility Comparison:

Crypto avg annualized vol: 11.06%
Equity avg annualized vol: 7.07%
Ratio (Crypto/Equity): 1.56x



7. Market Microstructure Analysis

```
In [18]: def analyze_microstructure(data, symbol):
    try:
        results = {}
        returns = data['returns'].dropna()
        acf_1 = returns.autocorr(lag=1)
        acf_5 = returns.autocorr(lag=5)
        acf_10 = returns.autocorr(lag=10)

        lb_result = acorr_ljungbox(returns, lags=10, return_df=True)
        lb_stat = lb_result['lb_stat'].iloc[-1]
        lb_pval = lb_result['lb_pvalue'].iloc[-1]
    except:
        lb_stat, lb_pval = None, None

    if 'volume' in data.columns:
        volume = data['volume'].dropna()
        volume_cv = volume.std() / volume.mean() if volume.mean() > 0 else None
        volume_autocorr = volume.autocorr(lag=1)
        price_volume_corr = data['returns'].abs().corr(data['volume'])
    else:
        volume_cv = None
        volume_autocorr = None
        price_volume_corr = None

    if 'high' in data.columns and 'low' in data.columns:
        hl_ratio = np.log(data['high'] / data['low'])
        beta = hl_ratio**2
        gamma = (np.log(data['high'] / data['low'].shift(1)))**2
        alpha = (np.sqrt(2*beta) - np.sqrt(beta)) / (3 - 2*np.sqrt(2)) - np.sqrt(beta)
        spread_cs = 2 * (np.exp(alpha) - 1) / (1 + np.exp(alpha))
        avg_spread = spread_cs.mean()
        spread_vol = spread_cs.std()
    else:
        avg_spread = None
        spread_vol = None

    if 'volume' in data.columns:
        dollar_volume = data['volume'] * data['close']
        amihud = (returns.abs() / (dollar_volume + 1e-10)).mean()
```

```

else:
    amihud = None

results = {
    'symbol': symbol,
    'autocorrelation': {
        'lag_1': acf_1,
        'lag_5': acf_5,
        'lag_10': acf_10
    },
    'ljung_box': {
        'statistic': lb_stat,
        'p_value': lb_pval,
        'has_serial_correlation': lb_pval < 0.05 if lb_pval else None
    },
    'volume': {
        'coefficient_variation': volume_cv,
        'autocorrelation': volume_autocorr,
        'price_volume_corr': price_volume_corr
    },
    'liquidity': {
        'avg_spread_cs': avg_spread,
        'spread_volatility': spread_vol,
        'amihud_illiquidity': amihud
    }
}

return results
except Exception as e:
    print(f"Error in microstructure analysis for {symbol}: {str(e)}")
    return None

print("MARKET MICROSTRUCTURE")
print("="*80)

print(f"Analyzing microstructure for {len(all_symbols)} symbols...\n")

microstructure_results = {}

for i, symbol in enumerate(all_symbols, 1):
    if symbol in analysis_data:
        print(f"[{i}/{len(all_symbols)}] Processing {symbol}...", end=' ')
        result = analyze_microstructure(analysis_data[symbol], symbol)
        if result:
            microstructure_results[symbol] = result
            print("Pass")
        else:
            print("Fail")

for symbol in all_symbols:
    if symbol in microstructure_results:
        micro_stats = microstructure_results[symbol]

        print(f"\n{symbol}:")
        print(f"Autocorrelation:")
        if micro_stats['autocorrelation']['lag_1'] is not None:

```

```

        print(f"Lag 1: {micro_stats['autocorrelation']['lag_1']:.4f}")
    if micro_stats['autocorrelation']['lag_5'] is not None:
        print(f"Lag 5: {micro_stats['autocorrelation']['lag_5']:.4f}")
    if micro_stats['autocorrelation']['lag_10'] is not None:
        print(f"Lag 10: {micro_stats['autocorrelation']['lag_10']:.4f}")

    if micro_stats['ljung_box']['p_value'] is not None:
        print(f"Ljung-Box test:")
        print(f"p-value: {micro_stats['ljung_box']['p_value']:.4f}")
        print(f"Serial correlation: {'Yes' if micro_stats['ljung_box']['has_ser']

    if micro_stats['volume']['coefficient_variation'] is not None:
        print(f"Volume patterns:")
        print(f"Coefficient of variation: {micro_stats['volume']['coefficient_v
    if micro_stats['volume']['autocorrelation'] is not None:
        print(f"Volume autocorrelation: {micro_stats['volume']['autocorrela
    if micro_stats['volume']['price_volume_corr'] is not None:
        print(f"|Return|-Volume correlation: {micro_stats['volume']['price_

    if micro_stats['liquidity']['avg_spread_cs'] is not None:
        print(f"    Liquidity metrics:")
        print(f"        Avg Corwin-Schultz spread: {micro_stats['liquidity']['avg_s
    if micro_stats['liquidity']['spread_volatility'] is not None:
        print(f"        Spread volatility: {micro_stats['liquidity']['spread_vo

    if micro_stats['liquidity']['amihud_illiquidity'] is not None:
        print(f"        Amihud illiquidity: {micro_stats['liquidity']['amihud_illiq

```

MARKET MICROSTRUCTURE

=====

Analyzing microstructure for 8 symbols...

[1/8] Processing BTCUSD... Pass
[2/8] Processing ETHUSD... Pass
[3/8] Processing SOLUSD... Pass
[4/8] Processing XRPUSD... Pass
[5/8] Processing ADAUSD... Pass
[6/8] Processing SPY... Pass
[7/8] Processing QQQ... Pass
[8/8] Processing IWM... Pass

BTCUSD:

Autocorrelation:

Lag 1: -0.1626

Lag 5: -0.0005

Lag 10: 0.0015

Ljung-Box test:

p-value: 0.0000

Serial correlation: No

Volume patterns:

Coefficient of variation: 2.0875

Volume autocorrelation: 0.6114

|Return|-Volume correlation: 0.4484

Liquidity metrics:

Avg Corwin-Schultz spread: -0.000004

Spread volatility: 0.001615

Amihud illiquidity: 8.99e-09

ETHUSD:

Autocorrelation:

Lag 1: -0.0844

Lag 5: -0.0031

Lag 10: 0.0010

Ljung-Box test:

p-value: 0.0000

Serial correlation: No

Volume patterns:

Coefficient of variation: 2.1441

Volume autocorrelation: 0.5558

|Return|-Volume correlation: 0.5021

Liquidity metrics:

Avg Corwin-Schultz spread: -0.000005

Spread volatility: 0.002256

Amihud illiquidity: 4.23e-07

SOLUSD:

Autocorrelation:

Lag 1: -0.0463

Lag 5: -0.0005

Lag 10: 0.0022

Ljung-Box test:

p-value: 0.0000

Serial correlation: No

Volume patterns:

Coefficient of variation: 2.2690
Volume autocorrelation: 0.5020
|Return|-Volume correlation: 0.4441
Liquidity metrics:
Avg Corwin-Schultz spread: -0.000015
Spread volatility: 0.003274
Amihud illiquidity: 1.82e-05

XRPUSD:
Autocorrelation:
Lag 1: -0.0579
Lag 5: -0.0076
Lag 10: -0.0016
Ljung-Box test:
p-value: 0.0000
Serial correlation: No
Volume patterns:
Coefficient of variation: 2.5756
Volume autocorrelation: 0.6231
|Return|-Volume correlation: 0.4958
Liquidity metrics:
Avg Corwin-Schultz spread: -0.000069
Spread volatility: 0.003124
Amihud illiquidity: 7.26e-01

ADAUSD:
Autocorrelation:
Lag 1: -0.0889
Lag 5: -0.0055
Lag 10: 0.0021
Ljung-Box test:
p-value: 0.0000
Serial correlation: No
Volume patterns:
Coefficient of variation: 2.6699
Volume autocorrelation: 0.4691
|Return|-Volume correlation: 0.4086
Liquidity metrics:
Avg Corwin-Schultz spread: -0.000052
Spread volatility: 0.003015
Amihud illiquidity: 3.65e-02

SPY:
Autocorrelation:
Lag 1: -0.0170
Lag 5: -0.0022
Lag 10: 0.0047
Ljung-Box test:
p-value: 0.0009
Serial correlation: Yes
Volume patterns:
Coefficient of variation: 2.2519
Volume autocorrelation: 0.3850
|Return|-Volume correlation: 0.1241
Liquidity metrics:
Avg Corwin-Schultz spread: -0.000065

Spread volatility: 0.001472
Amihud illiquidity: 2.13e-10

QQQ:

Autocorrelation:

Lag 1: -0.0153

Lag 5: -0.0031

Lag 10: 0.0020

Ljung-Box test:

p-value: 0.0197

Serial correlation: Yes

Volume patterns:

Coefficient of variation: 1.8319

Volume autocorrelation: 0.4914

|Return|-Volume correlation: 0.1384

Liquidity metrics:

Avg Corwin-Schultz spread: -0.000086

Spread volatility: 0.002147

Amihud illiquidity: 3.55e-10

IWM:

Autocorrelation:

Lag 1: -0.0393

Lag 5: -0.0001

Lag 10: 0.0018

Ljung-Box test:

p-value: 0.0000

Serial correlation: Yes

Volume patterns:

Coefficient of variation: 2.1146

Volume autocorrelation: 0.3762

|Return|-Volume correlation: 0.1608

Liquidity metrics:

Avg Corwin-Schultz spread: -0.000166

Spread volatility: 0.001604

Amihud illiquidity: 1.66e-09

8. Hurst & Long Memory

```
In [21]: def calculate_hurst_exponent(time_series, max_lag=100):
    try:
        ts = np.asarray(time_series)
        lags = range(2, min(max_lag, len(ts) // 2))
        tau = []
        for lag in lags:
            diff = ts[lag:] - ts[:-lag]
            tau_val = np.sqrt(np.std(diff))
            tau.append(float(tau_val))
        lags_array = np.array(list(lags))
        tau_array = np.array(tau)
        valid_mask = (tau_array > 0) & np.isfinite(tau_array)
        if valid_mask.sum() < 2:
            return np.nan
        poly = np.polyfit(np.log(lags_array[valid_mask]), np.log(tau_array[valid_ma
        hurst = poly[0] * 2.0
```



```

        return hurst
    except Exception as e:
        print(f"Error calculating Hurst exponent: {str(e)}")
        return np.nan

def calculate_hurst_for_symbol(symbol, df):
    try:
        results = {}

        # Calculate for prices
        price_series = df['close'].values
        results['price'] = calculate_hurst_exponent(price_series)

        # Calculate for returns
        returns_series = df['returns'].dropna().values
        results['returns'] = calculate_hurst_exponent(returns_series)

        # Calculate for volume if available
        if 'volume' in df.columns:
            volume_series = df['volume'].values
            results['volume'] = calculate_hurst_exponent(volume_series)
        else:
            results['volume'] = None

        return symbol, results
    except Exception as e:
        print(f"Error processing {symbol}: {str(e)}")
        return symbol, {}

hurst_results = {}

for i, symbol in enumerate(all_symbols, 1):
    if symbol in analysis_data:
        print(f"[{i}/{len(all_symbols)}] Computing Hurst exponent for {symbol}...",
              sym, hurst_values = calculate_hurst_for_symbol(symbol, analysis_data[symbol])
        if hurst_values:
            hurst_results[sym] = hurst_values
            print("Pass")
        else:
            print("Error")

# Display results with interpretation
print("\nResults:")
print("="*80)

for symbol in all_symbols:
    if symbol in hurst_results:
        h = hurst_results[symbol]
        print(f"\n{symbol}:")

        # Price series
        if h.get('price') is not None and not np.isnan(h['price']):
            interpretation = 'Trending' if h['price'] > 0.55 else 'Mean-reverting'
            print(f"  Price series: H = {h['price']:.3f} [{interpretation}]")

        # Returns

```

```

        if h.get('returns') is not None and not np.isnan(h['returns']):
            interpretation = 'Trending' if h['returns'] > 0.55 else 'Mean-reverting'
            print(f"    Returns:      H = {h['returns']:.3f} [{interpretation}]")

        # Volume
        if h.get('volume') is not None and not np.isnan(h['volume']):
            interpretation = 'Trending' if h['volume'] > 0.55 else 'Mean-reverting'
            print(f"    Volume:      H = {h['volume']:.3f} [{interpretation}]")

if len(available_crypto) > 0 and len(available_equity) > 0:
    crypto_hurst_returns = [hurst_results[s]['returns'] for s in available_crypto
                            if s in hurst_results and hurst_results[s].get('returns')
                            and not np.isnan(hurst_results[s]['returns'])]
    equity_hurst_returns = [hurst_results[s]['returns'] for s in available_equity
                            if s in hurst_results and hurst_results[s].get('returns')
                            and not np.isnan(hurst_results[s]['returns'])]

    if crypto_hurst_returns and equity_hurst_returns:
        print(f"\nHurst Exponent Comparison (Returns):")
        print(f"    Crypto average: {np.mean(crypto_hurst_returns):.3f}")
        print(f"    Equity average: {np.mean(equity_hurst_returns):.3f}")
        print(f"    Difference: {np.mean(crypto_hurst_returns) - np.mean(equity_hurst_returns):.3f}")

print("\nHurst exponent analysis complete")

```

[1/8] Computing Hurst exponent for BTCUSD... Pass
[2/8] Computing Hurst exponent for ETHUSD... Pass
[3/8] Computing Hurst exponent for SOLUSD... Pass
[4/8] Computing Hurst exponent for XRPUSD... Pass
[5/8] Computing Hurst exponent for ADAUSD... Pass
[6/8] Computing Hurst exponent for SPY... Pass
[7/8] Computing Hurst exponent for QQQ... Pass
[8/8] Computing Hurst exponent for IWM... Pass

Results:

=====

BTCUSD:

Price series: H = 0.477 [Random walk]
Returns: H = -0.000 [Mean-reverting]
Volume: H = 0.057 [Mean-reverting]

ETHUSD:

Price series: H = 0.490 [Random walk]
Returns: H = -0.000 [Mean-reverting]
Volume: H = 0.047 [Mean-reverting]

SOLUSD:

Price series: H = 0.484 [Random walk]
Returns: H = -0.000 [Mean-reverting]
Volume: H = 0.039 [Mean-reverting]

XRPUSD:

Price series: H = 0.480 [Random walk]
Returns: H = -0.001 [Mean-reverting]
Volume: H = 0.059 [Mean-reverting]

ADAUSD:

Price series: H = 0.487 [Random walk]
Returns: H = -0.000 [Mean-reverting]
Volume: H = 0.038 [Mean-reverting]

SPY:

Price series: H = 0.493 [Random walk]
Returns: H = -0.000 [Mean-reverting]
Volume: H = 0.054 [Mean-reverting]

QQQ:

Price series: H = 0.494 [Random walk]
Returns: H = -0.000 [Mean-reverting]
Volume: H = 0.067 [Mean-reverting]

IWM:

Price series: H = 0.484 [Random walk]
Returns: H = -0.000 [Mean-reverting]
Volume: H = 0.046 [Mean-reverting]

Hurst Exponent Comparison (Returns):

Crypto average: -0.000
Equity average: -0.000
Difference: -0.000

9. Regime Detection and Analysis

```
In [22]: def identify_market_regimes(data, symbol):
df = data.copy()
window = 24 * 7
df['rolling_vol'] = df['returns'].rolling(window=window).std()
df['rolling_mean'] = df['returns'].rolling(window=window).mean()
vol_median = df['rolling_vol'].median()
conditions = [
    (df['rolling_mean'] > 0) & (df['rolling_vol'] <= vol_median),
    (df['rolling_mean'] > 0) & (df['rolling_vol'] > vol_median),
    (df['rolling_mean'] <= 0) & (df['rolling_vol'] <= vol_median),
    (df['rolling_mean'] <= 0) & (df['rolling_vol'] > vol_median),
]
choices = ['Bull_LowVol', 'Bull_HighVol', 'Bear_LowVol', 'Bear_HighVol']
df['regime'] = pd.Series(
    np.select(conditions, choices, default='Undefined'),
    index=df.index
)
regime_stats = df.groupby('regime').agg({
    'returns': ['mean', 'std', 'count'],
    'volume': 'mean' if 'volume' in df.columns else lambda x: None
})
return df, regime_stats

if len(all_symbols) > 0:
    symbol = all_symbols[0]
    if symbol in analysis_data:
        print(f"\n" + "="*80)
        print(f"REGIME ANALYSIS FOR {symbol}")
        regime_data, regime_stats = identify_market_regimes(analysis_data[symbol],

        print("\nStatistics:")
        print(regime_stats)
        # Visualize regimes
        fig, axes = plt.subplots(2, 1, figsize=(14, 10), sharex=True)

        ax = axes[0]
        regime_colors = {
            'Bull_LowVol': 'green',
            'Bull_HighVol': 'lightgreen',
            'Bear_LowVol': 'orange',
            'Bear_HighVol': 'red',
            'Undefined': 'gray'
        }

        ax.plot(regime_data.index, regime_data['close'], linewidth=1, color='black'

        for regime, color in regime_colors.items():
            mask = regime_data['regime'] == regime
            if mask.any():
                ax.fill_between(regime_data.index,
```

```

regime_data['close'].min(),
regime_data['close'].max(),
where=mask, alpha=0.2, color=color, label=regime)

ax.set_title(f'{symbol} Price with Market Regimes', fontweight='bold')
ax.set_ylabel('Price')
ax.legend(loc='upper left')
ax.grid(True, alpha=0.3)

# Regime distribution over time
ax = axes[1]
regime_dummies = pd.get_dummies(regime_data['regime'])
regime_cumsum = regime_dummies.rolling(window=24*30).mean() # 30-day average

for col in regime_cumsum.columns:
    if col in regime_colors:
        ax.plot(regime_cumsum.index, regime_cumsum[col],
                label=col, color=regime_colors[col], linewidth=2)

ax.set_title('Regime Prevalence (30-Day Rolling Average)', fontweight='bold')
ax.set_ylabel('Proportion')
ax.set_xlabel('Date')
ax.legend()
ax.grid(True, alpha=0.3)


plt.tight_layout()
plt.show()

```

REGIME ANALYSIS FOR BTCUSD

Statistics:

	returns			volume
	mean	std	count	mean
regime				
Bear_HighVol	-0.000061	0.001085	328388	16.515578
Bear_LowVol	-0.000030	0.000433	327009	5.891170
Bull_HighVol	0.000062	0.001029	351092	16.560892
Bull_LowVol	0.000030	0.000441	352472	6.374267
Undefined	0.000002	0.000565	167	14.953849



10. Anomaly Detection

```
In [23]: def detect_anomalies(data, symbol, contamination=0.01):
features = ['returns', 'volume'] if 'volume' in data.columns else ['returns']
# Add technical features
df = data.copy()
df['returns_abs'] = df['returns'].abs()
df['returns_squared'] = df['returns'] ** 2
if 'high' in df.columns and 'low' in df.columns:
    df['high_low_ratio'] = df['high'] / df['low']
    features.extend(['high_low_ratio'])
features.extend(['returns_abs', 'returns_squared'])
# Prepare data
X = df[features].dropna()
iso_forest = IsolationForest(contamination=contamination, random_state=42)
anomalies = iso_forest.fit_predict(X)
# Add anomaly labels back to dataframe
df.loc[X.index, 'anomaly'] = anomalies
df['is_anomaly'] = df['anomaly'] == -1
return df

if len(all_symbols) > 0:
    symbol = all_symbols[0]
    if symbol in analysis_data:
        print(f"ANOMALY DETECTION - {symbol}")
        print("="*80)
        anomaly_data = detect_anomalies(analysis_data[symbol], symbol)
        n_anomalies = anomaly_data['is_anomaly'].sum()
```

```

pct_anomalies = n_anomalies / len(anomaly_data) * 100
print(f"\nTotal anomalies detected: {n_anomalies} ({pct_anomalies:.2f}%)")
normal_stats = anomaly_data[~anomaly_data['is_anomaly']]['returns'].describe()
anomaly_stats = anomaly_data[anomaly_data['is_anomaly']]['returns'].describe()
comparison = pd.DataFrame({
    'Normal': normal_stats,
    'Anomalous': anomaly_stats
})
print("\nReturn Statistics Comparison:")
print(comparison)
# Visualize anomalies
fig, axes = plt.subplots(2, 1, figsize=(14, 10))
ax = axes[0]
ax.plot(anomaly_data.index, anomaly_data['returns'],
        linewidth=0.5, color='blue', alpha=0.5, label='Returns')

# Highlight anomalies
anomaly_points = anomaly_data[anomaly_data['is_anomaly']]
ax.scatter(anomaly_points.index, anomaly_points['returns'],
          color='red', s=20, alpha=0.7, label='Anomalies')

ax.set_title(f'{symbol} Returns with Detected Anomalies', fontweight='bold')
ax.set_ylabel('Returns')
ax.legend()
ax.grid(True, alpha=0.3)

# Price with anomalies
ax = axes[1]
ax.plot(anomaly_data.index, anomaly_data['close'],
        linewidth=1, color='black', alpha=0.7)
# Mark anomaly periods
for idx in anomaly_points.index:
    ax.axvline(x=idx, color='red', alpha=0.1, linewidth=0.5)
ax.set_title(f'{symbol} Price with Anomaly Periods', fontweight='bold')
ax.set_ylabel('Price')
ax.set_xlabel('Date')
ax.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()

print("\nAnomaly detection complete")



```

ANOMALY DETECTION - BTCUSD

Total anomalies detected: 13586 (1.00%)

Return Statistics Comparison:

	Normal	Anomalous
count	1.345542e+06	13586.000000
mean	1.060392e-06	0.000069
std	6.670888e-04	0.004637
min	-3.650029e-03	-0.036938
25%	-2.822343e-04	-0.003795
50%	0.000000e+00	0.002209
75%	2.849229e-04	0.003779
max	3.584933e-03	0.036828



Anomaly detection complete

11. Feature Importance from Models

```
In [ ]: Not Implemented
```

12. Comprehensive Summary Statistics

```
In [26]: # Create comprehensive summary
summary_stats = {}

print("COMPREHENSIVE STATISTICAL SUMMARY")
```



```

for symbol in all_symbols:
    if symbol in analysis_data:
        df = analysis_data[symbol]
        returns = df['returns'].dropna()

        # Calculate comprehensive statistics
        summary_stats[symbol] = {
            # Basic statistics
            'count': len(returns),
            'mean_return': returns.mean(),
            'median_return': returns.median(),
            'std_return': returns.std(),

            # Risk metrics
            'sharpe_ratio': returns.mean() / returns.std() * np.sqrt(365 * 24) if r
            'sortino_ratio': returns.mean() / returns[returns < 0].std() * np.sqrt(
            'max_drawdown': (df['close'] / df['close'].cummax() - 1).min(),
            'var_95': np.percentile(returns, 5),
            'cvar_95': returns[returns <= np.percentile(returns, 5)].mean(),

            # Higher moments
            'skewness': skew(returns),
            'kurtosis': kurtosis(returns),
            'jarque_bera_pval': jarque_bera(returns)[1],

            # Market structure
            'hurst_exponent': hurst_results.get(symbol, {}).get('returns', None) if
            'autocorr_lag1': returns.autocorr(lag=1),

            # Volatility
            'annualized_vol': returns.std() * np.sqrt(365 * 24),
            'vol_of_vol': returns.rolling(window=24).std().std(),
        }

# Create summary DataFrame
summary_df = pd.DataFrame(summary_stats).T

# Display key metrics
print("\nKey Risk-Adjusted Performance Metrics:")
print(summary_df[['sharpe_ratio', 'sortino_ratio', 'max_drawdown', 'annualized_vol']

print("\nDistribution Characteristics:")
print(summary_df[['skewness', 'kurtosis', 'jarque_bera_pval']].round(4))

print("\nMarket Microstructure:")
if 'hurst_exponent' in summary_df.columns:
    print(summary_df[['hurst_exponent', 'autocorr_lag1']].round(4))
else:
    print(summary_df[['autocorr_lag1']].round(4))

if len(available_crypto) > 0 and len(available_equity) > 0:
    print("\n" + "="*80)
    print("CRYPTO vs EQUITY COMPARISON")
    print("="*80)

    crypto_stats = summary_df.loc[available_crypto].mean()

```

```
equity_stats = summary_df.loc[available_equity].mean() if len(available_equity)

if len(equity_stats) > 0:
    comparison = pd.DataFrame({
        'Crypto (Avg)': crypto_stats,
        'Equity (Avg)': equity_stats,
        'Difference': crypto_stats - equity_stats
    })

    print(comparison.round(4))
else:
    print("\n[INFO] No equity data available for comparison")
    print("\nCrypto Statistics (Average):")
    print(crypto_stats.round(4))
```

COMPREHENSIVE STATISTICAL SUMMARY

Key Risk-Adjusted Performance Metrics:

	sharpe_ratio	sortino_ratio	max_drawdown	annualized_vol
BTCUSD	0.2016	0.2487	-0.3312	0.0758
ETHUSD	0.1231	0.1508	-0.6617	0.0866
SOLUSD	0.2032	0.2679	-0.6757	0.1311
XRPUSD	0.1775	0.2103	-0.5871	0.1268
ADAUSD	0.1186	0.1464	-0.6565	0.1326
SPY	0.4422	0.8906	-0.0895	0.0569
QQQ	0.5693	1.3164	-0.0901	0.0884
IWM	0.2079	0.3292	-0.1589	0.0667

Distribution Characteristics:

	skewness	kurtosis	jarque_bera_pval
BTCUSD	-0.2911	44.6983	0.0
ETHUSD	-2.1143	293.5962	0.0
SOLUSD	1.7546	301.6893	0.0
XRPUSD	-0.6628	270.5865	0.0
ADAUSD	0.8059	542.0845	0.0
SPY	135.7958	29773.7545	0.0
QQQ	171.2137	40704.7271	0.0
IWM	54.5236	8408.1725	0.0

Market Microstructure:

	hurst_exponent	autocorr_lag1
BTCUSD	-0.0002	-0.1626
ETHUSD	-0.0001	-0.0844
SOLUSD	-0.0004	-0.0463
XRPUSD	-0.0009	-0.0579
ADAUSD	-0.0003	-0.0889
SPY	-0.0003	-0.0170
QQQ	-0.0001	-0.0153
IWM	-0.0004	-0.0393

CRYPTO vs EQUITY COMPARISON

	Crypto (Avg)	Equity (Avg)	Difference
count	1.346073e+06	74170.3333	1.271903e+06
mean_return	0.000000e+00	0.0000	-0.000000e+00
median_return	0.000000e+00	0.0000	0.000000e+00
std_return	1.200000e-03	0.0008	4.000000e-04
sharpe_ratio	1.648000e-01	0.4065	-2.417000e-01
sortino_ratio	2.048000e-01	0.8454	-6.406000e-01
max_drawdown	-5.824000e-01	-0.1128	-4.696000e-01
var_95	-1.600000e-03	-0.0007	-1.000000e-03
cvar_95	-2.700000e-03	-0.0011	-1.600000e-03
skewness	-1.015000e-01	120.5110	-1.206125e+02
kurtosis	2.905310e+02	26295.5514	-2.600502e+04
jarque_bera_pval	0.000000e+00	0.0000	0.000000e+00
hurst_exponent	-4.000000e-04	-0.0003	-1.000000e-04
autocorr_lag1	-8.800000e-02	-0.0239	-6.410000e-02
annualized_vol	1.106000e-01	0.0707	3.990000e-02
vol_of_vol	7.000000e-04	0.0006	1.000000e-04

Save Results for Next Notebook

```
In [25]: results_to_save = {
    'analysis_timestamp': datetime.now().isoformat(),
    'symbols_analyzed': all_symbols,
    'crypto_symbols': available_crypto,
    'equity_symbols': available_equity,
    'analysis_data': analysis_data,
    'distribution_results': distribution_results,
    'correlation_matrix': correlation_matrix if 'correlation_matrix' in locals() else None,
    'volatility_results': volatility_results if 'volatility_results' in locals() else None,
    'microstructure_results': microstructure_results if 'microstructure_results' in locals() else None,
    'hurst_results': hurst_results if 'hurst_results' in locals() else None,
    'summary_statistics': summary_df if 'summary_df' in locals() else None,
    'engineered_data': engineered_data,
    'model_results': model_results,
    'config': config
}

# Save to pickle file
output_file = 'notebooks/02_statistical_results.pkl'
with open(output_file, 'wb') as f:
    pickle.dump(results_to_save, f)
```